

Siddhi Galada

(818) 272-0721 | Los Angeles, CA | galada@usc.edu | [linkedin.com/in/siddhi-galada](https://www.linkedin.com/in/siddhi-galada) | github.com/SiddhiGalada44 | [Portfolio](#)

EDUCATION

University of Southern California

M.S. in Computer Science — Coursework: Algorithms, NLP, Machine Learning, Information Retrieval

Los Angeles, CA
Aug 2024 – May 2026 (Expected)

Sir MVIT

B.E. in Computer Science and Engineering, CGPA: 9.09/10

Bengaluru, India

Dec 2020 – Jul 2024

TECHNICAL SKILLS

Languages: Java, Go (Golang), Python, C++, JavaScript, TypeScript, SQL, C

Frameworks & Tools: Node.js, Express.js, Spring Boot, Flask, React, REST APIs, Microservices, JWT, OAuth2, RBAC, Git, Docker

AI & ML: LLMs (Anthropic Claude, OpenAI GPT), Prompt Engineering, RAG Pipelines, Vector Databases (FAISS, ChromaDB), Embeddings & Semantic Search, AI Agents & Tool Use, Fine-tuning (DistilBERT, BERT), LangChain

Cloud/DevOps: Google Cloud (GKE, Cloud Run, Firestore, GCR), AWS (EC2, EKS), Docker, Kubernetes, CI/CD (GitHub Actions), Linux

Databases: PostgreSQL, MongoDB, MySQL, Firestore

EXPERIENCE

Suki.ai

Software Engineering Intern

Redwood City, CA

Jun 2025 – Aug 2025

- Architected SimEHR, a FHIR R4-compliant simulated EHR platform exposing REST APIs to replace unstable vendor sandboxes and enable deterministic end-to-end integration testing for the Suki Assistant pipeline.
- Designed stateless Go microservices for routing, validation, and audit logging across 10+ workflows with structured logging, retry middleware, and error handling to improve fault tolerance and observability.
- Containerized and deployed to Google Cloud (GKE, Cloud Run) via Docker and Kubernetes, cutting environment setup from 4–5 hours to under 5 minutes and improving deployment reproducibility by 95%.
- Built automated REST API compatibility tests in CI/CD (GitHub Actions), detecting integration regressions early and improving pre-production service stability.

Varcons Technologies Pvt Ltd

Software Engineering Intern – NLP Systems

Bengaluru, India

Aug 2023 – Oct 2023

- Built a Python ETL pipeline processing 500,000+ tweets with a Flask REST API for prediction endpoints; containerized with Docker and deployed to AWS.
- Integrated a fine-tuned BERT model achieving 89% accuracy; authored 30+ unit tests covering 80% of endpoint logic, reducing post-deployment defects across edge cases.

PROJECTS

PromptLens – LLM Prompt Evaluation Framework [\[Link\]](#)

Python, Anthropic Claude API, OpenAI API, LangChain, FAISS, FastAPI, React

Los Angeles, CA

Jan 2026 – Mar 2026

- Built a prompt evaluation framework benchmarking LLM variants (Claude, GPT-4o) on faithfulness and hallucination rate; RAG pipeline with FAISS and cosine similarity retrieval reduced hallucination by 38% over naive prompting baselines.
- Designed chain-of-thought and few-shot templating with structured output parsing; exposed results via FastAPI backend and React dashboard surfacing per-prompt latency, token cost, and quality scores.

Distributed Key-Value Store [\[Link\]](#)

Python, gRPC, Raft Consensus, Docker

Los Angeles, CA

Sep 2025 – Jan 2026

- Built a fault-tolerant, linearizable key-value store implementing Raft consensus for leader election, log replication, and safe failover across a 5-node cluster with a gRPC-based client-server communication layer.
- Implemented log compaction via snapshotting to reduce follower catch-up time by 60%; validated with chaos testing under randomized node crashes and network partitions.

Shared Contacts Service – Tenant-aware Backend [\[Link\]](#)

Java, Spring Boot, PostgreSQL, JWT, Docker

Los Angeles, CA

Sep 2025 – Oct 2025

- Built a tenant-aware REST service isolating data across 5+ tenants via X-Tenant-ID interceptor and @Version optimistic locking, eliminating write conflicts under concurrent load.
- Configured JWT auth and RBAC (Admin, Member) with Spring Security; paginated endpoints sustained 100+ concurrent requests at sub-100ms P95 latency, packaged with Docker Compose.

GitWise – Code Search and Repository Navigation Platform [\[Link\]](#)

TypeScript, Node.js, Express.js, React, Docker

Los Angeles, CA

Aug 2025 – Oct 2025

- Developed a full-stack platform for exploring large codebases; implemented a repository indexing pipeline supporting 5,000+ files with sub-150ms query latency via metadata extraction and in-memory caching.
- Designed low-latency REST endpoints handling 50+ concurrent requests and reduced environment setup from 20 minutes to under 2 minutes via Docker Compose orchestration.